

Punctuation Sensitivity Benchmark

Comparative analysis of F0 and pause behaviour across four TTS models
— Chatterbox, XTTS-v2, F5-TTS, and Kokoro —
across 28 utterances spanning 5 punctuation categories.

1 Executive Summary

This benchmark evaluates how four TTS models interpret punctuation in prosody. The central finding is a **clear trade-off between F0-based and pause-based punctuation sensitivity**:

- Chatterbox** leads in F0 modulation: question marks trigger a **+68.5 Hz/s** rising slope versus periods, and quotation shifts the F0 range by 50 Hz.
- XTTS-v2** excels at pause-driven hierarchy: its comma-to-ellipsis pause gradient (101.5 ms) is the largest among all models, yielding a perfect hierarchy score of 1.0.
- F5-TTS** shows moderate F0 sensitivity (question rise +31.8 Hz/s) but has the weakest pause hierarchy (0.33).
- Kokoro** produces minimal F0 variation across punctuation and serves as an energetic baseline with moderate pause differentiation.

Key Results Summary

Model	Question Rise (Hz/s)	Pause Hierarchy	Quotation	Best At
Chatterbox	+68.5	0.67	YES	F0 cues
XTTS-v2	-5.0	1.00	NO	Pause timing
F5-TTS	+31.8	0.33	YES	—
Kokoro	+25.7	0.67	NO	baseline

2 Methodology

Stimuli. 28 unique utterances were manually designed to probe each punctuation category. Each utterance was rendered with the exact punctuation required (no textual variation beyond the target mark) to isolate the effect of punctuation on prosody.

Categories (5). The utterances span *sentence-end* (period, question mark, exclamation mark), *pause hierarchy* (comma, semicolon, em-dash, ellipsis), *trailing punctuation* (ellipsis vs period), *quotation* (quoted speech vs reported speech), and *capitalisation* (ALL CAPS, Title Case, lowercase).

Models (4). Chatterbox, XTTS-v2, F5-TTS, and Kokoro were each used to generate 28 audio clips at 24 kHz. Synthesis used default sampling parameters (temperature=0.7, top-k=40 where applicable) with no voice cloning

or speaker conditioning.

Acoustic analysis. Voice activity was detected using an energy-based VAD (50 ms window, 20 ms step, relative threshold 0.15). F0 was extracted with `pyin` (50–600 Hz range, 10 ms step). Pause durations were measured as silences ≥ 30 ms between voiced regions. Terminal F0 slope was computed over the final 100 ms of each voiced region via linear regression.

Gate check. A sanity gate was applied before analysis: each model's period-marked utterances were verified to have at least 50% of frames voiced in their final segment. All four models passed every gate item.


3 Gate Check Results

All models passed the sanity gate with a 100% pass rate on period-ending utterances.

Model	Period Items	Passing	Pass Rate	Status
Chatterbox	5	5	100%	PASS
Kokoro	5	5	100%	PASS
XTTS-v2	5	5	100%	PASS
F5-TTS	5	5	100%	PASS

4 Sentence-End Sensitivity

The figure below shows mean terminal F0 slope for each punctuation type. Positive slopes indicate rising pitch (question-like); negative slopes indicate falling pitch (statement-like). The key metric is the **question-vs-period difference**: a model that genuinely rises for questions should show a large positive difference.



Model	Period (Hz/s)	Question (Hz/s)	Exclamation (Hz/s)	Q-P Diff (Hz/s)
Chatterbox	-8.3	60.2	-106.5	68.5
Kokoro	-251.9	-226.2	-488.5	25.7
XTTS-v2	-83.0	-88.0	-158.9	-5.0
F5-TTS	-87.0	-55.2	-489.0	31.8

Interpretation. Chatterbox shows the strongest question-mark differentiation (+68.5 Hz/s relative to periods), followed by F5-TSS (+31.8 Hz/s) and Kokoro (+25.7 Hz/s). XTTS-v2 is essentially flat (-5.0 Hz/s), treating questions and statements with nearly identical F0 trajectories. For exclamation marks, F5-TTS produces the steepest downward slope (-489 Hz/s), indicating an aggressive terminal fall.

5 Pause Hierarchy

Punctuation marks imply different pause lengths: commas are brief, ellipses are long. The hierarchy score measures how consistently a model orders pauses as comma < semicolon < em-dash < ellipsis. A score of 1.0 means perfect ordering across all pairs.




Model	Comma (ms)	Semicolon (ms)	Em-dash (ms)	Ellipsis (ms)	Hierarchy Score
Chatterbox	112.5	119.1	110.8	118.1	0.7
Kokoro	135.8	154.5	126.9	166.7	0.7
XTTS-v2	146.2	194.6	220.0	247.6	1.0
F5-TTS	218.8	354.7	341.3	334.5	0.3

Interpretation. XTTS-v2 achieves a perfect hierarchy score (1.0), with a clear monotonic increase from commas (146 ms) to ellipses (248 ms). Chatterbox and Kokoro tie at 0.67 – both show the correct general trend but with occasional inversions (e.g., em-dash shorter than semicolon). F5-TTS scores lowest (0.33); its pause durations are dominated by a high baseline – all marks produce pauses ≥ 200 ms – and the ordering between categories is inconsistent.

6 Trailing Punctuation: Ellipsis vs Period

Trailing punctuation (ellipsis vs period) was compared on F0 slope, amplitude decay, and pause duration. A model sensitive to the trailing mark should show different F0 and pause behaviour between the two.




Model	Ellipsis F0 (Hz/s)	Period F0 (Hz/s)	Ellipsis Pause (ms)	Period Pause (ms)	F0 Diff (Hz/s)
Chatterbox	-124.6	-148.8	235.0	220.0	24.2
Kokoro	63.1	2.4	500.0	505.0	60.7
XTTS-v2	-60.9	-43.8	590.0	690.0	-17.1
F5-TTS	-86.0	-87.2	1285.0	1305.0	1.1

Interpretation. Kokoro shows the largest F0 difference between ellipsis and period (+60.7 Hz/s), indicating distinct prosodic treatment of suspension vs finality. Chatterbox also differentiates (+24.2 Hz/s). XTTS-v2 and F5-TTS show negligible differences (-17.1 and +1.1 Hz/s respectively), suggesting they do not prosodically distinguish trailing ellipsis from a period. Pause durations are broadly similar within each model across the two conditions.

7 Quotation Sensitivity

Quoted speech was compared to reported (non-quoted) speech. The metrics are F0 mean shift and F0 range shift between the two conditions. A shift-detected flag indicates whether the model changes its F0 profile for quotation marks.





Model	F0 Mean Shift (Hz)	F0 Range Shift (Hz)	Shift Detected
Chatterbox	8.5	50.0	YES
Kokoro	5.8	3.7	NO
XTTS-v2	1.8	3.9	NO
F5-TTS	12.4	11.7	YES

Interpretation. Chatterbox shows the strongest quotation sensitivity, with a 50 Hz F0 range expansion — the widest shift of any model — clearly separating quoted from reported speech. F5-TTS also shifts F0 range (+11.7 Hz) and the shift is flagged as detected. XTTS-v2 and Kokoro produce negligible range shifts (3.9 Hz and 3.7 Hz) and were not flagged, indicating they do not modulate F0 in response to quotation marks.

8 Summary: The F0 vs Pause Trade-Off

The scatter plot below visualises the central finding: models cluster along a diagonal trade-off between F0-based sensitivity (question-vs-period rise) and pause-based sensitivity (comma-to-ellipsis pause gradation).



Model	Question Rise (Hz/s)	Pause Gradient (ms)	Dominant Strategy
Chatterbox	68.5	5.6	F0-driven
Kokoro	25.7	30.8	Mixed
XTTS-v2	-5.0	101.5	Pause-driven
F5-TTS	31.8	115.8	Pause-driven

Key finding. Models that rely on F0 modulation for punctuation sensitivity (Chatterbox, F5-TTS) show weaker pause hierarchies, while the model with the strongest pause hierarchy (XTTS-v2) shows no F0-based punctuation differentiation. This trade-off suggests two distinct prosodic strategies in current TTS architectures.

9 Capitalisation Sensitivity

All-caps vs title-case vs lowercase utterances were compared for RMS amplitude boost and F0 boost relative to the lowercase baseline.

Model	RMS Boost in ALL CAPS	F0 Boost in ALL CAPS
Chatterbox	YES	NO
Kokoro	YES	—
XTTS-v2	YES	0.0
F5-TTS	YES	YES

Interpretation. All models increase RMS amplitude for ALL CAPS text, suggesting a universal "louder = emphatic" strategy. F0 boost in ALL CAPS is less common: only F5-TTS shows a clear F0 elevation. Chatterbox and XTTS-v2 do not raise F0 for all-caps, while Kokoro's data was inconclusive.

10 Limitations

- **Sample size.** 28 utterances per model (112 total) yield robust descriptive statistics but limited inferential power. Category-level N ranges from 2 to 13, and some conditions (e.g., quotation pairs) have only 2–4 items per model.
- **Single-speaker.** All audio was generated in a single synthetic voice per model. Results may not generalise to voice-cloned or multi-speaker scenarios.
- **VAD tuning.** Energy-based VAD parameters (threshold 0.15, 50 ms window) were optimised for this dataset; different settings may shift absolute pause durations, though relative comparisons should be robust.
- **F0 tracking.** `pyin` is robust but can produce spurious values in unvoiced or creaky regions. We required $\geq 50\%$ voiced frames in terminal segments for gate inclusion.
- **Limited punctuation scope.** We tested five categories; other marks (colon, parentheses, hyphenation) were not included.
- **No listening test.** These are purely acoustic metrics. Perceptual naturalness and appropriateness of prosody were not evaluated.
- **Single temperature.** The default sampling temperature (0.7) may conceal stochastic variation in punctuation sensitivity.